An Elaboration on

Relative Entropy Under Mappings by Stochastic Matrices
Joel E. Cohen, Yoh Iwasa, Gh. Rautu, Mary Beth Ruskai, Eugene Seneta, Gh. Zbaganu

By Aaron Kelley

1. Introduction

The properties of relative entropy have been extensively studied in the past in many different contexts and under a variety of different names, but its properties under mappings by stochastic matrices is the main focus of this paper. The results can be used to provide more information on bounding the rates of convergence to equilibrium of ergodic Markov chains and Markov processes. After reviewing some definitions, we will see theorem 3.1 (although omiting the proof) which will be used to prove the the two main results selected for this elaboration: theorems 4.1 and 5.4. Finally, we will see the some of the ideas of this paper in the context of the what we saw throughout the semester.

2. Background

Preliminary Definitions

Definitions 1.1 (Vectors): Let m, n, and d be finite positive integers. Vectors that are $n \times 1$ or $d \times 1$ will be called **n** and **d** vectors respectively. We define

$$N_d = \{x \in R^d : x_i \geq 0, \sum_i x_i = 1\} \quad P_d = \{x \in N_d : x_i > 0, \ \forall i\}$$

Definitions 1.2 (Matrices): A **(column) stochastic** $m \times n$ matrix is a matrix whos columns belong to N_m. A nonnegative matrix is called **row-allowable** if each row contains at least one positive element. A matrix with at least one positive row (all elements of a row positive) is called **row-positive**. A column-stochastic row-positive matrix is called a **Markov matrix**. A nonnegative $d \times d$ matrix A is called **primitive** if A^k is positive for some positive integer k. A column-stochastic $m \times n$ matrix is called a **scrambling** matrix if any submatrix consisting of two columns has a row both elements of which are positive. Note, every row-positive matrix is scrambling, but not conversely.

Main Definitions

Definition 2.1 (Symmetric Relative Entropy): For any two positive d-vectors $x = (x_i)$ and $y = (y_i)$, whether or not x and y are probability vectors, we define the *relative entropy* as

$$H(x, y) = \sum_i x_i \log(x_i / y_i)$$

and they *symmetric relative entropy* as

$$J(x,y) = H(x,y) + H(y,x) = \sum_i (x_i - y_i) \log \frac{x_i}{y_i}$$

Definition 2.2 (relative ϕ-entropy): Let ϕ be a continuous real-valued function on $(0,\infty)\times(0,\infty)$ that is homogeneous and jointly convex in its arguments, and satisfies $\phi(1,1) = 0$. For any two positive d-vectors, $x = (x_i), y = (y_i)$, whether or not x and y are probability vectores, we defind the *relativeϕ-entropy* as

$$H_\phi = \sum_i \phi(x_i, y_i)$$

and the *symmetric relative ϕ-entropy* as

$$J_\phi(x,y) = H_\phi(x,y) + H_\phi(y,x)$$

Because $\tilde{\phi}(a,b) = \phi(a,b) + \phi(b,a)$ satisfies the conditions of 2 if ϕ does, and $J_\phi(x,y) = H_{\tilde{\phi}}(x,y)$, from this point on we will speak of J_ϕ as $H_{\tilde{\phi}}$

The function ϕ defined in definition 2 is jointly convex in both arguments if and only if $g(t) \equiv \phi(1, 1+t)$ is convex for $t \in (-1,\infty)$. Therefore any continuous real-valued convex function $g(t)$ on $(-1,\infty)$ such that $g(0) = 0$ defines a relative ϕ-entropy via the assumptions that $\phi(1,1+t) = g(t)$ and ϕ is homogeneous. So the relative ϕ-entropy and related quantities can be indexed by both ϕ and/or g. That's to say

$$H_\phi(x,y) = \sum_i \phi(x_i, y_i) \iff H_g(x,y) = \sum_i x_i g(y_i/x_i - 1)$$

Keep in mind that in all cases $H_{\log}$ denotes the relative entropy in **??**. That is

$$H_{\log} = H_g, \quad \text{when } g(t) = -\log(1+t)$$

Three main properties of relative entropy:

1. H_ϕ is a continuous, real-valued function that is homogeneous, jointly convex in (x,y) for any positive d-vectors x and y, subadditive, and such that $H_\phi(x,x) = 0$.

2. For any $x, y \in P_d$, $H_\phi(x,y) \geq 0$; and if $\phi(1,t)$ is strictly convex for $t \in (0,\infty)$, then $H_\phi(x,y) = 0$ if and only if $x = y$.

3. For any positive d-vectors x, y and positive n-vectors x, y any permutation matrices Q_1, Q_2 of size $m \times m$ and $n \times n$, respectively, and any row-allowable $m \times n$ matrix A, there exists positive n-vectors x', y' such that

$$\frac{H_\phi(Q_1 A Q_2 x, Q_1 A Q_2 y)}{H_\phi(x,y)} = \frac{H_\phi(Ax', Ay')}{H_\phi(x',y')}$$

4. If A is a column-stochastic, row-allowable $m \times d$ matrix and x, y are positive d-vectors, and $\phi(1,\cdot)$ convex, then $H_\phi(Ax, Ay) \leq H_\phi(x,y)$

Definition 2.3 (Dobrushin's Ergodicity Coefficient): For any $m \times n$ matrix A, Dobrushin's coefficient of ergodicity is

$$\alpha(A) = \min_{j,k} \sum_{i=1}^{m} \min(a_{ij}, a_{ik})$$

We will see that the complement, $1 - \alpha(A)$, is a bit more interesting with respect to the conclusions that we arrive at. Here we note that

$$\bar{\alpha}(A) \equiv 1 - \alpha(A) = \frac{1}{2} \max_{j,k} \sum_{i=1}^{m} |a_{ij} - a_{ik}|$$

and also satisfies

$$\bar{\alpha}(A) = \sup \left\{ \frac{\|A(x-y)\|_1}{\|x-y\|_1} : \ x \text{ and } y \text{ are positive } n\text{-vectors such that } x \neq y, \text{ y } \|x\|_1 = \|y\|_1 \right\}$$

Definition 2.4 (ϕ-entropy contraction coefficient): Let A be a column-stochastic, row-allowable $m \times n$ matrix. We define the **relative ϕ-enropy contraction coefficient**

$$\eta_\phi(A) = \sup \left\{ \frac{H_\phi(Ax, Ay)}{H_\phi(x, y)} : x, y \in P_n, x \neq y \right\}$$

3. Principal Results

For this section, assume that $\phi(1, \cdot)$ is strictly convex on $(0, \infty)$, $x, y \in P_n$, $x \neq y$.
Theorem 3.1: Let A be a column-stochastic, row-allowable $m \times n$ matrix, and let $x, y \in P_n$. Then

$$H_\phi(Ax, Ay) \leq \bar{\alpha}(A) H_\phi(x, y)$$

Theorem 4.1: $0 \leq \eta_\phi(A) \leq \bar{\alpha}(A) \leq 1$
Proof: Given theorem 3.1, it's trivial. Under the assumption of strict convexity of $\phi(1, \cdot)$, by property 2 we have that $0 < H_\phi(x, y)$ and as A is a nonnegative matrix, Ax and Ay are nonnegative (possibly equal) d-vectors, so $0 \leq H_\phi(Ax, Ay)$. Clearly

$$0 \leq H_\phi(Ax, Ay), \ 0 < H_\phi(x, y) \implies 0 \leq \left\{ \frac{H_\phi(Ax, Ay)}{H_\phi(x, y)} \right\}$$

From Theorem 3.1 (**ref this**), we have

$$\frac{H_\phi(Ax, Ay)}{H_\phi(x, y)} \leq \bar{\alpha}(A) \implies \sup \left\{ \frac{H_\phi(Ax, Ay)}{H_\phi(x, y)} \right\} \leq \bar{\alpha}(A)$$

$\bar{\alpha}(A) \leq 1$ follows directly from property 4 (**ref this**) of relative entropy. Thus

$$0 \leq \sup\left\{\frac{H_\phi(Ax, Ay)}{H_\phi(x, y)}\right\} \leq \bar{\alpha}(A) \leq 1$$

□

Theorem 5.4: If $g(w)$ is thrice differentiable in a neighborhood of 0 and $g''(0) > 0$, then $\eta_{w^2}(A) \leq \eta_g(A)$; in particular, $\eta_{w^2}(A) \leq \eta_{\log}(A)$
Proof:

$$\begin{aligned}\eta_{\omega^2}(A) &= \sup_{\substack{x \neq y \\ x,y \in P_n}} \frac{H_{\omega^2}(Ax, Ay)}{H_{\omega^2}(x, y)} \\ &= \sup_{\substack{x \in P_n \\ v^\top 1 = 0}} \frac{H_{\omega^2}(Ax, Ax + Av)}{H_{\omega^2}(x, x + v)} \equiv \sup_{\substack{x \in P_n \\ v^\top 1 = 0}} \frac{\Phi(Ax, Av)}{\Phi(x, v)} \end{aligned} \tag{1}$$

Where $\Phi(x, v) = H_{\omega^2}(x, x + v) = \sum_j \frac{v_j^2}{x_j}$ and $v \neq 0$. Now we use the fact that g is thrice differentiable to expand it in a Taylor's series about $w = 0$.

$$\phi(s, t) = sg\left(\frac{t}{s} - 1\right) = (t - s)g'(0) + \frac{(t - s)^2}{s}\frac{g''(0)}{2} + \frac{1}{s^2}O\left((t - s)^3\right)$$

Pick ϵ sufficiently small and let $y_\epsilon = x + \epsilon v$. Then

$$H_g(x, y_\epsilon) = \frac{g''(0)}{2}\epsilon^2\Phi(x, v) + O(\epsilon^3)$$

and

$$H_g(Ax, Ay_\epsilon) = \frac{g''(0)}{2}\epsilon^2\Phi(Ax, Av) + O(\epsilon^3)$$

therefore,

$$\eta_g(A) \geq \frac{H_g(Ax, Ay_\epsilon)}{H_g(x, y_\epsilon)} = \frac{\Phi(Ax, Av)}{\Phi(x, v)} + O(\epsilon)$$

Because of equation (1), we can choose x, v and ϵ such that $\frac{\Phi(Ax,Av)}{\Phi(x,v)} + O(\epsilon)$ is arbitrarily close to $n_{\omega^2}(A)$. It follows that $\eta_{w^2}(A) \leq \eta_g(A)$ □

4. Relation to Course

In the course, we first defined and studied the properties of KL-divergence (in fact, this is the relative entropy defined in Definition 2.1, also called $H_{\log}$ throughout the paper), and generalized it to a family of divergences called f-divergence. This paper also forms a generalization of KL-divergence, but instead derives a more abstract version of relative entropy. These two methods of generalization can be connected, and we'll see that relative entropy is a type of f-divergence.

Let $f : (0, \infty) \to \mathbb{R}$ be convex and $f(1) = 0$. Furthermore, let P and Q be probability

distributions such that $P << Q$. Define f-divergence to be

$$D_f(P||Q) := \int f\left(\frac{dP}{dQ}\right) dQ \xrightarrow{\text{discrete}} \sum_i f\left(\frac{p_i}{q_i}\right) q_i$$

where $p = (p_i)$ and $q = (q_i)$ are two d-vectors. Thus f-divergence (in the discrete case) is equivalente to the relative g-entropy if and only if $f(z) = zg\left(\frac{1}{z} - 1\right)$ [2]. Then

$$D_f(P||Q) = \sum_i f\left(\frac{p_i}{q_i}\right) q_i = \sum_i \frac{p_i}{\cancel{q_i}} \cancel{q_i} g\left(\frac{1}{p_i/q_i}\right) = \sum_i p_i g\left(\frac{q_i}{p_i} - 1\right) = H_g(P, Q)$$

slightly abusing notation when we think about P and Q as vectors of their respective probabilities indexed at i. For $z \in (0, \infty)$, $\frac{1}{z} - 1 \in (-1, \infty)$ so we have that $z \mapsto zg\left(\frac{1}{z} - 1\right)$ is convex. And of course, $f(1) = 1g(1-1) = g(0) = 0$; thus, everything is well defined.

On a related note, we can also extend f divergence to be symmetric in a similar way to how the symmetric relative entropy is defined, $J_f(x, y) = D_f(x, y) + D_f(y, x)$. Note that this still does not permit us to think about divergence as a metric, as it still does not satisfy the triangle inequality.

Lastly, we can draw some parallels about the Data Processing Inequality that we've seen in class. A is a scrambling matrix if and only if $\eta_\phi(A) < 1$; furthermore, we have that $n_\phi(A) = 1$ if A is a permutation matrix. Clearly, by definition of $\eta_\phi(A)$ we have that $H_\phi(Ax, Ay) \leq H_\phi(x, y)$ with equality if and only if A is a permutation matrix. In class we saw the DPI in the context of divergence,

$$D(P_{Y|X}P_X||P_{Y|X}Q_X) \leq D(P_x||Q_x)$$

Because of how the contraction coefficient was defined in this paper, we can easily set $P_{X|Y} = A$ and $P_{Y|Z} = A$ and set up the markov chain such that

$$X \xrightarrow{A} Y \xrightarrow{A'} Z \implies H_\phi(X, Y) \leq H_\phi(X, Z)$$

where the equality only holds if both A and A' are permutation matrices.

5. Conclusion

We can think of relative entropy from Q to P (or KL Divergence) as the *information gain*. That is, we have a random variable X we are using Q as it's distribution. The relative entropy tells us how much information we could gain about X if we were to use P instead of Q. To expand this idea into a more abstract sense, we need to define a continuous homogeneous function $\phi(x, y)$ convex in both arguments satisfying $\phi(1, 1) = 0$. Then the calculation is simple, $\sum_i \phi(x_i, y_i)$. Using $g(t) \equiv \phi(1, 1 + t)$ to index the relative ϕ-entropy facilitated many of the proofs and lines of thought presented in this paper. Defining and elaborating the contraction coefficient $\eta_g(A)$ gives many usefull inequalities in studing the effects of the (stochastic) matrix A on probability vectores x, y through the lense of relative entropy.

[2] Note that g is not necessarily invertible